

# Mileage Prediction - Regression Analysis

## Source:

This dataset was taken from the StatLib library which is maintained at Carnegie Mellon University. The dataset was used in the 1983 American Statistical Association Exposition.

## Data Set Information:

This dataset is a slightly modified version of the dataset provided in the StatLib library. In line with the use by Ross Quinlan (1983) in predicting the attribute "mpg", 8 of the original instances were removed because they had unknown values for the "mpg" attribute. the original dataset is available in the file "auto-mpg.data-original". "The data concerned city-cycle fuel composition in miles per gallon, to be predicted in terms of 3 multivalued discrete and 5 continuous attributes."(Quinlan, 1993)

## Attribute Information:

- 1.mpg: continuous
- 2.cylinders: multivalued discrete
- 3.displacement: continuous
- 4.horsepower: continuous
- 5.weight: continuous
- 6.acceleration: continuous
- 7.model year: multivalued discrete
- 8.origin: multivalued discrete
- 9.car name: string(unique for each instance)

## ✓ Import Library

Double-click (or enter) to edit

```
import pandas as pd
```

```
import numpy as np
```

```
import matplotlib.pyplot as plt
```

```
import seaborn as sns
```

## ✓ Import Data

```
df = pd.read_csv('https://github.com/YBI-Foundation/Dataset/raw/main/MPG.csv')
```

```
df.head()
```



	mpg	cylinders	displacement	horsepower	weight	acceleration	model_year	origin
0	18.0	8	307.0	130.0	3504	12.0	70	usa
1	15.0	8	350.0	165.0	3693	11.5	70	usa
2	18.0	8	318.0	150.0	3436	11.0	70	usa
3	16.0	8	304.0	150.0	3433	12.0	70	usa
4	17.0	8	302.0	140.0	3449	10.5	70	usa



Next steps:

[Generate code with df](#)

[View recommended plots](#)

```
df.nunique()
```



```
mpg          129
cylinders      5
displacement  82
horsepower    93
weight       351
acceleration  95
model_year    13
origin        3
name         305
dtype: int64
```

## ✓ Data Preprocessing

```
df.info()
```



```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 398 entries, 0 to 397
Data columns (total 9 columns):
 #   Column          Non-Null Count  Dtype
---  -
```

```

-----
0  mpg          398 non-null  float64
1  cylinders    398 non-null  int64
2  displacement 398 non-null  float64
3  horsepower   392 non-null  float64
4  weight       398 non-null  int64
5  acceleration 398 non-null  float64
6  model_year   398 non-null  int64
7  origin       398 non-null  object
8  name         398 non-null  object
dtypes: float64(4), int64(3), object(2)
memory usage: 28.1+ KB

```

```
df.describe()
```



	mpg	cylinders	displacement	horsepower	weight	acceleration	model_year
<b>count</b>	398.000000	398.000000	398.000000	392.000000	398.000000	398.000000	398.000000
<b>mean</b>	23.514573	5.454774	193.425879	104.469388	2970.424623	15.568090	7.000000
<b>std</b>	7.815984	1.701004	104.269838	38.491160	846.841774	2.757689	0.000000
<b>min</b>	9.000000	3.000000	68.000000	46.000000	1613.000000	8.000000	7.000000
<b>25%</b>	17.500000	4.000000	104.250000	75.000000	2223.750000	13.825000	7.000000
<b>50%</b>	23.000000	4.000000	148.500000	93.500000	2803.500000	15.500000	7.000000
<b>75%</b>	29.000000	8.000000	262.000000	126.000000	3608.000000	17.175000	7.000000
<b>max</b>	46.600000	8.000000	455.000000	230.000000	5140.000000	24.800000	8.000000

```
df.corr()
```



```

-----
ValueError                                Traceback (most recent call last)
<ipython-input-115-2f6f6606aa2c> in <cell line: 1>()
----> 1 df.corr()

-----
3 frames -----
/usr/local/lib/python3.10/dist-packages/pandas/core/internals/managers.py in
_interleave(self, dtype, na_value)
    1792         else:
    1793             arr = blk.get_values(dtype)
-> 1794         result[r1.indexer] = arr
    1795         itemmask[r1.indexer] = 1
    1796

ValueError: could not convert string to float: 'usa'

```

Next steps:

[Explain error](#)

## ✓ Remove Missing Value


```
df = df.dropna()
```

```
df.info()
```


```
<class 'pandas.core.frame.DataFrame'>
Index: 392 entries, 0 to 397
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype  
---  -
0   mpg              392 non-null   float64
1   cylinders        392 non-null   int64  
2   displacement     392 non-null   float64
3   horsepower       392 non-null   float64
4   weight           392 non-null   int64  
5   acceleration     392 non-null   float64
6   model_year      392 non-null   int64  
7   origin           392 non-null   object  
8   name             392 non-null   object  
dtypes: float64(4), int64(3), object(2)
memory usage: 30.6+ KB
```

## ✓ Data Visualization

```
sns.pairplot(df, x_vars= ['displacement', 'horsepower', 'weight', 'acceleration', 'mpg'], y='mpg')
```



```
sns.regplot(x = 'displacement', y = 'mpg', data = df);
```



## ✓ Define Target Variable y and Feature X

```
df.columns
```

```
Index(['mpg', 'cylinders', 'displacement', 'horsepower', 'weight',  
      'acceleration', 'model_year', 'origin', 'name'],  
      dtype='object')
```

```
y = df['mpg']
```

```
y.shape
```

```
(392,)
```

```
X = df[['displacement', 'horsepower', 'weight', 'acceleration']]
```

```
X.shape
```

```
(392, 4)
```

```
X
```



	displacement	horsepower	weight	acceleration	
0	307.0	130.0	3504	12.0	
1	350.0	165.0	3693	11.5	
2	318.0	150.0	3436	11.0	
3	304.0	150.0	3433	12.0	
4	302.0	140.0	3449	10.5	
...	...	...	...	...	
393	140.0	86.0	2790	15.6	
394	97.0	52.0	2130	24.6	
395	135.0	84.0	2295	11.6	
396	120.0	79.0	2625	18.6	
397	119.0	82.0	2720	19.4	

392 rows × 4 columns

Next steps:

[Generate code with X](#)[View recommended plots](#)

## ✓ Scaling Data

```
from sklearn.preprocessing import StandardScaler
```

```
ss = StandardScaler()
```

```
X = ss.fit_transform(X)
```

```
X
```

```
array([[ 1.07728956,  0.66413273,  0.62054034, -1.285258  ],
       [ 1.48873169,  1.57459447,  0.84333403, -1.46672362],
       [ 1.1825422 ,  1.18439658,  0.54038176, -1.64818924],
       ...,
       [-0.56847897, -0.53247413, -0.80463202, -1.4304305 ],
       [-0.7120053 , -0.66254009, -0.41562716,  1.11008813],
       [-0.72157372, -0.58450051, -0.30364091,  1.40043312]])
```

```
pd.DataFrame(X).describe()
```



	0	1	2	3
<b>count</b>	3.920000e+02	3.920000e+02	3.920000e+02	3.920000e+02
<b>mean</b>	-7.250436e-17	-1.812609e-16	-1.812609e-17	4.350262e-16
<b>std</b>	1.001278e+00	1.001278e+00	1.001278e+00	1.001278e+00
<b>min</b>	-1.209563e+00	-1.520975e+00	-1.608575e+00	-2.736983e+00
<b>25%</b>	-8.555316e-01	-7.665929e-01	-8.868535e-01	-6.410551e-01
<b>50%</b>	-4.153842e-01	-2.853488e-01	-2.052109e-01	-1.499869e-02
<b>75%</b>	7.782764e-01	5.600800e-01	7.510927e-01	5.384714e-01
<b>max</b>	2.493416e+00	3.265452e+00	2.549061e+00	3.360262e+00



After Standardization Mean is Zero and Standard Deviation is One

## ✓ Train Test Split Data

```
from sklearn.model_selection import train_test_split
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, train_size = 0.7, random_state =
```

```
X_train.shape, X_test.shape, y_train.shape, y_test.shape
```

```
((274, 4), (118, 4), (274,), (118,))
```



## ✓ Linear Regression Model

```
from sklearn.linear_model import LinearRegression
```

```
lr = LinearRegression()
```

```
lr.fit(X_train, y_train)
```



```
LinearRegression()
LinearRegression()
```

```
lr.intercept_
```



```
23.485738559737584
```

Mileage = 23.4 = 1.05Displacement- 1.68Horsepower- 4.10Weight- 0.115Acceleration + error

## ✓ Predict Test Data

```
y_pred = lr.predict(X_test)
```

```
y_pred
```

```
→ array([18.51865637, 15.09305675, 14.30128789, 23.6753321 , 29.7546115 ,
        23.68796629, 26.61066644, 24.56692437, 15.06260986, 11.94312046,
        24.08050053, 27.96518468, 31.66130278, 31.01309132, 18.32428976,
        19.32795009, 28.08847536, 32.1506879 , 31.15859692, 27.15792144,
        18.82433097, 22.54580176, 26.15598115, 32.36393869, 20.74377679,
         8.78027518, 22.19699435, 18.20614294, 25.00052718, 15.26421552,
        23.13441082, 17.10542257,  9.87180062, 30.00790415, 20.41204655,
        29.11860245, 24.4305187 , 21.72601835, 10.51174626, 13.12426391,
        21.41938406, 19.96113872,  6.19146626, 17.79025345, 22.5493033 ,
        29.34765021, 13.4861847 , 25.88852083, 29.40406946, 22.41841964,
        22.07684766, 16.46575802, 24.06290693, 30.12890046, 10.11318121,
         9.85011438, 28.07543852, 23.41426617, 20.08501128, 30.68234133,
        20.92026393, 26.78370281, 22.9078744 , 14.15936872, 24.6439883 ,
        26.95515832, 15.25709393, 24.11272087, 30.80980589, 14.9770217 ,
        27.67836372, 24.2372919 , 10.92177228, 30.22858779, 30.88687365,
        27.33992044, 31.18447082, 10.8873597 , 27.63510608, 16.49231363,
        25.63229888, 29.49776285, 14.90393439, 32.78670687, 30.37325244,
        30.9262743 , 14.71702373, 27.09633246, 26.69933806, 29.06424799,
        32.45810182, 29.44846898, 31.61239999, 31.57891837, 21.46542321,
        31.76739191, 26.28605476, 28.96419915, 31.09628395, 24.80549594,
        18.76490961, 23.28043777, 23.04466919, 22.14143162, 15.95854367,
        28.62870918, 25.58809869, 11.4040908 , 25.73334842, 30.83500051,
        21.94176255, 15.34532941, 30.37399213, 28.7620624 , 29.3639931 ,
        29.10476703, 20.44662365, 28.11466839])
```

## ✓ Model Accuracy

```
from sklearn.metrics import mean_absolute_error, mean_absolute_percentage_error, r2_score
```

```
mean_absolute_error(y_test, y_pred)
```

```
→ 3.3286968643244106
```

```
mean_absolute_percentage_error(y_test, y_pred)
```

```
→ 0.14713035779536746
```

## ✓ Polynomial Regression



```
from sklearn.preprocessing import PolynomialFeatures
```

```
poly = PolynomialFeatures(degree=2, interaction_only = True, include_bias = False)
```

```
X_train2 = poly.fit_transform(X_train)
```

```
X_test2 = poly.fit_transform(X_test)
```

```
lr.fit(X_train2, y_train)
```



```
▼ LinearRegression  
LinearRegression()
```

```
lr.intercept_
```



```
21.27336450063766
```

```
lr.coef_
```



```
array([-2.76070596, -5.00559628, -1.36884133, -0.81225214,  1.24596571,  
       -0.12475017, -0.90542822,  1.35064048, -0.17337823,  1.41680398])
```

```
y_pred_poly = lr.predict(X_test2)
```

## ✓ Model Accuracy

```
from sklearn.metrics import mean_absolute_error, mean_absolute_percentage_error, r2_score
```

```
mean_absolute_error(y_test, y_pred_poly)
```



```
2.7887147720295977
```

```
mean_absolute_percentage_error(y_test, y_pred_poly)
```



```
0.12074018342938687
```

```
r2_score(y_test, y_pred_poly)
```